

1. Explain the elements of dynamic programming with an example.

method used in computer science and mathematics to solve problems by breaking them down into simpler sub problems. It is particularly useful for optimization problems where the same sub problems are solved multiple times. mainly an optimization over plain recursion. Wherever we see a recursive solution that has repeated calls for the same inputs,

Elements of Dynamic Programming:

1. Optimal Substructure

- A problem exhibits optimal substructure if its solution can be built from solutions to its subproblems.
- Example: The shortest path problem in a graph—if the shortest path from A to C goes through B, then the path from A to B and B to C must also be the shortest.

2. Overlapping Subproblems

- The problem can be broken down into smaller subproblems, and these subproblems are solved multiple times.
- Example: Fibonacci sequence—calculating Fibonacci(5) requires Fibonacci(4) and Fibonacci(3), but Fibonacci(4) also requires Fibonacci(3), leading to redundant calculations.

3. Memoization (Top-Down Approach)

- Stores the results of expensive function calls and returns the cached result when the same inputs occur again.
- Example: Storing already computed Fibonacci numbers in an array to avoid recomputation.

4. Tabulation (Bottom-Up Approach)

- Solves the problem iteratively by filling up a table, starting from the smallest subproblem.
- Example: Computing Fibonacci numbers iteratively from Fibonacci(0) to Fibonacci(n) and storing them in an array.

Example: Fibonacci Sequence Using DP

Let's compute Fibonacci numbers using both approaches.

1. Memoization (Top-Down)

Algorithms and Complexity Theory

python

Copy Edit

```
def fib_memo(n, memo={}):
    if n <= 1:
        return n
    if n not in memo:
        memo[n] = fib_memo(n-1, memo) + fib_memo(n-2, memo)
    return memo[n]

print(fib_memo(10)) # output: 55
```

Here, previously computed values are stored in `memo` to avoid redundant calculations.

2. Tabulation (Bottom-Up)

python

Copy Edit

```
def fib_tab(n):
    if n <= 1:
        return n
    dp = [0] * (n + 1)
    dp[1] = 1
    for i in range(2, n + 1):
        dp[i] = dp[i-1] + dp[i-2]
    return dp[n]

print(fib_tab(10)) # output: 55
```

Here, we build the solution iteratively using an array `dp`.

Dynamic Programming helps solve complex problems efficiently by using previously computed results, avoiding redundant calculations. It is widely used in problems like **Knapsack, Shortest Path, Matrix Chain Multiplication, Longest Common Subsequence, and more.**

1.

2. Explain the elements of greedy strategy with an example.

Elements of the Greedy Strategy

The **greedy strategy** is an algorithmic paradigm used to solve optimization problems by making a sequence of choices. Each choice is made based on a **local optimum**, with the hope that these local decisions lead to a **globally optimal solution**.

Key Elements of the Greedy Strategy

1. Greedy Choice Property

- A **locally optimal choice** at each step leads to a globally optimal solution.
- Example: In coin change (using denominations like 1, 5, 10), always picking the largest coin available is optimal in many cases.

2. Optimal Substructure

- The **optimal solution to a problem contains optimal solutions to its subproblems**.
- Example: In the shortest path problem (Dijkstra's algorithm), the shortest path to a node remains optimal even when extended to the next node.

3. Feasibility Check

- The chosen solution must remain valid after making a greedy choice.

4. Solution Construction

- The algorithm **iteratively builds the solution** by adding locally optimal choices.
-

Example: Activity Selection Problem

Problem Statement

Given **n activities** with their start and end times, select the **maximum number of non-overlapping activities**.

Greedy Approach

1. **Sort activities by their end time.**
2. **Select the first activity** (earliest finishing).
3. **Iterate through remaining activities** and pick the next one that starts **after the last selected activity ends**.

Implementation

```
python
def activity_selection(activities):
    # Sort activities based on end time
    activities.sort(key=lambda x: x[1])

    selected = []
    last_end_time = 0

    for start, end in activities:
        if start >= last_end_time:
            selected.append((start, end))
            last_end_time = end    # Update end time

    return selected

# Example Activities (start, end)
activities = [(1, 3), (2, 5), (3, 9), (6, 8), (5, 7)]
optimal_selection = activity_selection(activities)

print("Selected Activities:", optimal_selection)
```

Time Complexity: $O(n \log n)$ (due to sorting)

Conclusion

- **Greedy algorithms** work well when the **greedy choice property and optimal substructure hold**.
- **Limitations:** They do **not always guarantee an optimal solution** (e.g., in the **Knapsack problem**, a greedy strategy may fail).
- **Common Applications:** Huffman coding, Kruskal's algorithm, Dijkstra's algorithm, Prim's algorithm.

Algorithms and Complexity Theory

3. What is an optimal Huffman code for the following set of frequencies, based on the first 8 Fibonacci numbers?

a:1 b:1 c:2 d:3 e:5 f:8 g:13 h:21

Generalize your answer to find the optimal code when the frequencies are the first n Fibonacci numbers?

Optimal Huffman Code for Given Fibonacci Frequencies

Huffman coding is a greedy algorithm used for lossless data compression. It assigns shorter binary codes to more frequent characters and longer codes to less frequent ones.

Step 1: Given Frequencies

We are given the first 8 Fibonacci numbers as symbol frequencies:

Symbol Frequency

a	1
b	1
c	2
d	3
e	5
f	8
g	13
h	21

Step 2: Building the Huffman Tree

1. Combine the two smallest frequencies (1 and 1 → a and b) → New node with frequency 2.
2. Combine the next two smallest nodes (2 and 2 → (a,b) and c) → New node with frequency 4.
3. Combine the next two smallest nodes (3 and 4 → d and (a,b,c)) → New node with frequency 7.
4. Combine the next two smallest nodes (5 and 7 → e and (a,b,c,d)) → New node with frequency 12.
5. Combine the next two smallest nodes (8 and 12 → f and (a,b,c,d,e)) → New node with frequency 20.

6. **Combine the next two smallest nodes** (13 and 20 → g and (a,b,c,d,e,f)) → New node with frequency 33.
7. **Combine the last two nodes** (21 and 33 → h and (a,b,c,d,e,f,g)) → Final root node with frequency 54.

Step 3: Assigning Huffman Codes

We assign 0 to one branch and 1 to the other recursively:

Symbol Huffman Code

h	0
g	10
f	110
e	1110
d	11110
c	111110
a	1111110
b	1111111

Generalization for the First n Fibonacci Numbers

For the first n Fibonacci numbers, the Huffman tree follows a **greedy hierarchical merging**:

1. The two smallest nodes (both 1s) are combined first.
2. The tree grows by always merging the smallest available Fibonacci numbers.
3. The resulting pattern is **left-heavy**, meaning the largest frequency gets the shortest code.

Code Pattern:

- The highest frequency (last Fibonacci number) gets the shortest code (0).
- The second-highest gets one more bit (10).
- The third-highest gets 110, and so on.
- The lowest frequency items get the longest codes (111...).

This structure ensures an **optimal prefix code** while maintaining minimal encoding length for frequent symbols.

4. Suppose we perform a sequence of n operations on a data structure in which the i th operation costs i if i is an exact power of 2, and 1 otherwise. Use aggregate analysis to determine the amortized cost per operation.

Amortized Cost Analysis Using Aggregate Analysis

Given Cost Function

We perform n operations on a data structure where:

- If i (the operation index) is an exact power of 2, the cost is i .
- Otherwise, the cost is 1.

Step 1: Total Cost Calculation (Aggregate Analysis)

The total cost $T(n)$ is the sum of:

1. Costs of operations that are exact powers of 2.
2. Costs of all other operations, each contributing 1.

Identifying Costs of Power-of-2 Operations

The power-of-2 operations occur at:

$$i = 1, 2, 4, 8, 16, \dots$$

Algorithms and Complexity Theory

The total number of such operations is approximately $\log_2 n$ (since the largest power of 2 $\leq n$ is $2^{\lfloor \log_2 n \rfloor}$).

The contribution of these power-of-2 operations to the total cost is:

$$1 + 2 + 4 + 8 + \dots + 2^{\lfloor \log_2 n \rfloor}$$

Since this forms a geometric series:

$$\sum_{k=0}^{\lfloor \log_2 n \rfloor} 2^k = 2^{\lfloor \log_2 n \rfloor + 1} - 1$$

Approximating:

$$\sum_{k=0}^{\log_2 n} 2^k \approx 2n - 1 \approx 2n$$

Identifying Costs of Non-Power-of-2 Operations

For all other operations, the cost is simply **1 per operation**. Since there are $n - \log_2 n$ such operations, their total cost is:

$$n - \log_2 n$$

Total Cost $T(n)$

Adding both contributions:

$$T(n) = 2n + (n - \log_2 n) = 3n - \log_2 n$$

Step 2: Amortized Cost Per Operation

By aggregate analysis, the amortized cost per operation is:

$$\frac{T(n)}{n} = \frac{3n - \log_2 n}{n} = 3 - \frac{\log_2 n}{n}$$

For large n , the term $\frac{\log_2 n}{n}$ approaches 0, so the amortized cost is:

3

Final Answer

The **amortized cost per operation** is **O(3)=O(1)**. This shows that, despite occasional expensive operations (powers of 2), the average cost per operation remains constant.

Would you like further clarifications or a different approach?

5. Analyse and explain the randomized quick algorithm for its best case, average case and worst case time complexity.

Randomized QuickSort: Analysis & Complexity

Randomized QuickSort is a variation of the QuickSort algorithm where we randomly select the pivot at each step instead of always choosing a fixed position (like the first or last element). This randomization helps mitigate the worst-case scenario of QuickSort, making it efficient on average.

Algorithm Steps

1. Choose a random pivot from the array.
2. Partition the array into two subarrays:
 - Left subarray (elements smaller than the pivot).
 - Right subarray (elements greater than the pivot).
3. Recursively sort the two subarrays.
4. Combine the results to get the sorted array.

Time Complexity Analysis

1. Best Case ($O(n \log n)$)

- The best case occurs when the pivot divides the array into two equal halves at every step.
- If the partitioning is perfectly balanced, each recursive step splits the problem into two subproblems of size $n/2$, $n/4$, and so on.
- The recurrence relation for this case is:

$$T(n) = 2T(n/2) + O(n)$$

Solving using the recurrence tree method:

- Level 0: n operations
- Level 1: $2(n/2) = n$
- Level 2: $4(n/4) = n$
- ...
- Level $\log_2 n$: $n/n = 1$

Since there are $O(\log n)$ levels and each level takes $O(n)$ work, the total time is:

$$O(n \log n)$$

2. Average Case ($O(n \log n)$)

- The average case occurs when the pivot **partitions the array reasonably well but not perfectly**.
- On average, the pivot **divides the array into two subarrays of approximately $3n/4$ and $n/4$** .
- The recurrence relation:

$$T(n) = T(3n/4) + T(n/4) + O(n)$$

This still leads to $O(n \log n)$ complexity through recurrence expansion.

Summary of Time Complexities

Case	Complexity	Condition
Best Case	$O(n \log n)$	When pivot always splits array evenly.
Average Case	$O(n \log n)$	Random pivot gives roughly balanced partitions.
Worst Case	$O(n^2)$	Pivot consistently picks min/max element (bad partitioning).

$$T(n) = T(n - 1) + O(n)$$

Expanding:

$$\begin{aligned} T(n) &= T(n - 1) + n \\ &= T(n - 2) + (n - 1) + n \\ &= \dots = O(n^2) \end{aligned}$$

- This occurs in **deterministic QuickSort** when the array is **already sorted** and the pivot is always chosen as the first or last element.
- Randomization mitigates this** by ensuring that no consistent bad pivot choice occurs.



Why Randomization Works?

- Ensures that the pivot is not always bad** (avoids worst-case consistently).
- Probability distribution of pivot selection** keeps partitions balanced in expectation.
- Leads to expected $O(n \log n)$ time complexity in most cases**.

Conclusion

Randomized QuickSort is **one of the fastest sorting algorithms** in practice because of its expected $O(n \log n)$ time complexity. It is widely used due to its **low overhead and cache efficiency** compared to other $O(n \log n)$ sorts like Merge Sort.

- 6. Illustrate that performing a search in an n -node α -balanced amortized binary search tree takes $O(\log n)$ worst-case time.**

Search Complexity in an α -Balanced Amortized Binary Search Tree

1. Definition of an α -Balanced BST

An α -balanced binary search tree maintains a balance condition ensuring that for every subtree T of size n :

$$|T_{\text{left}}| \leq \alpha |T|$$

$$|T_{\text{right}}| \leq (1 - \alpha) |T|$$

where $0.5 \leq \alpha < 1$. This constraint ensures that the tree remains approximately balanced, preventing extreme cases like a linked list.

Common examples:

- **AVL Tree:** $\alpha \approx 0.5$ (ensures logarithmic height).
- **Red-Black Tree:** $\alpha = 2/3$, which allows slightly more imbalance but still maintains logarithmic height.

2. Height Analysis of an α -Balanced BST

The worst-case height $h(n)$ follows the recurrence:

$$h(n) = h(\alpha n) + 1$$

Expanding the recurrence:

$$h(n) = h(\alpha^2 n) + 2$$

$$h(n) = h(\alpha^3 n) + 3$$

Continuing until the base case $h(1) = 0$, we solve for k when $\alpha^k n = 1$:

$$k = \log_{1/\alpha} n$$

Thus, the height of the tree is:

$$h(n) = O(\log n)$$

Since a search operation in a BST requires traversing at most $h(n)$ levels, the worst-case search time is:

$$O(\log n)$$

3. Intuition Behind $O(\log n)$ Complexity

- The key idea is that **each level of recursion reduces the problem size by a constant factor α** .
 - Unlike an unbalanced BST (where height could be $O(n)$), the α -balanced property ensures that each subtree is **not too large**, leading to logarithmic depth.
 - **Self-balancing** mechanisms (such as AVL rotations or Red-Black Tree color changes) maintain the $O(\log n)$ height dynamically.
-

Conclusion

Since an α -balanced BST ensures that the height remains $O(\log n)$, a search operation—requiring at most $h(n)$ comparisons—runs in $O(\log n)$ **worst-case time**.

Graph and Geometric Algorithms

1. What are different types of graph representations used in graph algorithms? Explain with examples.

In graph algorithms, graphs can be represented in multiple ways, depending on the requirements of the algorithm. The three most common representations are:

1. Adjacency Matrix

An adjacency matrix is a 2D array (matrix) of size $V \times V$, where V is the number of vertices in the graph. Each cell $\text{matrix}[i][j]$ stores:

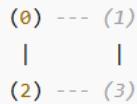
- 1 (or weight in weighted graphs) if there is an edge between vertex i and vertex j .
- 0 if there is no edge.

Example:

Consider an undirected graph:

lua

Copy Edit



Its adjacency matrix representation:

Copy Edit

```
0 1 1 0
1 0 0 1
1 0 0 1
0 1 1 0
```



Pros:

- Easy to check if an edge exists: $O(1)$ time complexity.
- Good for dense graphs.

Cons:

- Uses $O(V^2)$ space, which is inefficient for sparse graphs.

2. Adjacency List

An adjacency list uses an array of lists, where each index represents a vertex, and the corresponding list contains its adjacent vertices.

Example:

For the same graph:

```
css
```

Copy Edit

```
0 → [1, 2]
1 → [0, 3]
2 → [0, 3]
3 → [1, 2]
```

Pros:

- Uses $O(V + E)$ space, which is efficient for sparse graphs.
- Efficient for traversing all neighbors.

Cons:

- Checking if an edge exists takes $O(V)$ time in the worst case.

3. Edge List

An edge list is a simple representation where edges are stored as pairs (u, v) , and if weighted, as triplets (u, v, w) .

Example:

For the same graph:

```
css
```

Copy Edit

```
[(0,1), (0,2), (1,3), (2,3)]
```

Pros:

- Compact for sparse graphs.
- Useful for algorithms like Kruskal's.

Cons:

- Checking for an edge presence takes $O(E)$ time.
- Not efficient for adjacency queries.

Choosing the Right Representation:

- **Adjacency Matrix:** Suitable for dense graphs or when frequent edge lookups are needed.
- **Adjacency List:** Best for sparse graphs and efficient traversal.
- **Edge List:** Useful for algorithms that work on edges (e.g., Kruskal's algorithm).

2. Explain Push-relabel method for computing maximum flows.

Push-Relabel Algorithm for Maximum Flow

The **Push-Relabel method** (also known as the **Preflow-Push algorithm**) is an efficient algorithm for computing the **maximum flow** in a flow network. Unlike the Ford-Fulkerson method, which relies on augmenting paths, Push-Relabel maintains a **preflow** and works by redistributing flow using **push** and **relabel** operations.

Key Concepts

1. Preflow:

- Unlike standard flow algorithms, each vertex can hold excess flow (flow entering a vertex may exceed flow leaving it).
- Only the **source** and **sink** are allowed to have nonzero excess in the final solution.

2. Height Function (Labeling):

- Each vertex v is assigned a **height** $h(v)$ (or label), which ensures that flow moves downhill.
- The **source** is initially set to $h(s)=|V|$, while all other vertices start at height 0.

3. Push Operation:

- If a vertex has excess flow, it attempts to push flow to its neighbors.
- Flow can only be pushed from a higher-height vertex to a lower-height vertex.

4. Relabel Operation:

- If a vertex has excess flow but no available lower-height neighbors, it **increases** its height to enable flow movement.

Algorithm Steps

1. Initialization:

- Set the **height** of the source s to $|V|$.
- Set all flow in the edges from s to adjacent vertices to their capacity.
- Initialize **excess flow** in these vertices.

2. Loop until no more active vertices (vertices with excess flow other than source and sink):

- If a vertex has **excess flow**, try to **push** the flow to a lower-height neighbor.
- If no valid push is possible, **relabel** the vertex (increase its height).

3. Termination:

- When no more **active** vertices remain, the flow reaching the sink is the maximum flow.

Example

Consider a flow network:

Consider a flow network:



Step-by-Step Execution

1. Initialization:

- $h(s) = 4, h(A) = 0, h(B) = 0, h(t) = 0$
- Preflow pushes from s to A (5) and B (3).
- Excess flow: $e(A) = 5, e(B) = 3$.

2. Push operations:

- A pushes 4 to t (valid because $h(A) > h(t)$).
- B pushes 3 to t .

3. Maximum Flow:

- Total flow reaching t is $4 + 3 = 7$.

Time Complexity

- The Push-Relabel method runs in $O(V^2E)$ in the worst case.
- With the **FIFO** or **highest-label** optimizations, it can achieve $O(V^3)$, which is better for dense graphs.

Advantages

- More efficient than augmenting path methods like Ford-Fulkerson for dense graphs.
- Works well with optimizations (e.g., global relabeling).

3. Calculate the running time of BFS if we represent its input graph by an adjacency matrix and write a modified algorithm to handle this form of input?

Running Time of BFS Using an Adjacency Matrix

Breadth-First Search (BFS) explores all neighbors of a vertex before moving to the next level. The running time of BFS depends on how the graph is represented.

Time Complexity Analysis

- Let V be the number of vertices and E be the number of edges.
- With an **adjacency matrix**, checking for neighbors takes $O(V)$ time for each vertex.
- BFS processes each vertex **once** and checks all V possible neighbors.
- The total time complexity is:

$$O(V^2)$$

This is because:

- Visiting each vertex takes $O(V)$.
- Checking all V neighbors for each vertex results in $O(V \times V) = O(V^2)$.

Modified BFS Algorithm for an Adjacency Matrix

Since the adjacency matrix is a $V \times V$ matrix, BFS must check each row to find adjacent vertices.

Algorithm

```

python

from collections import deque

def bfs_adj_matrix(matrix, start):
    V = len(matrix) # Number of vertices
    visited = [False] * V
    queue = deque([start])
    visited[start] = True

    while queue:
        vertex = queue.popleft()
        print(vertex, end=" ") # Process the node

        # Check all possible neighbors
        for neighbor in range(V):
            if matrix[vertex][neighbor] == 1 and not visited[neighbor]:
                queue.append(neighbor)
                visited[neighbor] = True

# Example adjacency matrix (undirected graph)
graph = [
    [0, 1, 1, 0], # Edges from node 0
    [1, 0, 0, 1], # Edges from node 1
    [1, 0, 0, 1], # Edges from node 2
    [0, 1, 1, 0] # Edges from node 3
]

print("BFS traversal starting from node 0:")
bfs_adj_matrix(graph, 0)

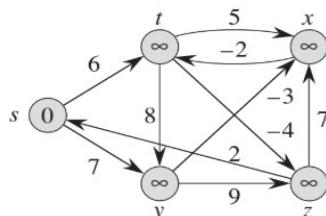
```

Comparison with Adjacency List

Representation	Space Complexity	BFS Time Complexity
Adjacency Matrix	$O(V^2)$	$O(V^2)$
Adjacency List	$O(V + E)$	$O(V + E)$

- Adjacency Matrix is inefficient for sparse graphs (many zeros).
- Adjacency List is better for graphs with fewer edges.

4. Run the Bellman-Ford algorithm on the directed graph of the Figure below , using vertex 's' as the source. In each pass, relax edges in the same order as in the figure, and show the d and π values after each pass.



Step-by-Step Execution of Bellman-Ford Algorithm

We will apply the **Bellman-Ford algorithm** on the directed graph using vertex s as the source. The **edge relaxation order** follows the structure of the graph.

Step 1: Initialization

For each vertex v , initialize:

- $d(s) = 0$ (distance from source to itself)
- $d(v) = \infty$ for all other vertices
- $\pi(v) = \text{NIL}$ (predecessor array to track paths)

Algorithms and Complexity Theory

Vertex	d (Distance)	π (Predecessor)
s	0	NIL
t	∞	NIL
x	∞	NIL
y	∞	NIL
z	∞	NIL

Step 2: Relaxation Process: Bellman-Ford relaxes all edges $|V|-1$ times.

Iteration 1 (Relax all edges in order)

Relax (s, t) with weight 6:

$$d(t) = \min(\infty, d(s) + 6) = 6, \pi(t) = s$$

Relax (s, y) with weight 7:

$$d(y) = \min(\infty, d(s) + 7) = 7, \pi(y) = s$$

Relax (t, x) with weight 5:

$$d(x) = \min(\infty, d(t) + 5) = 6 + 5 = 11, \pi(x) = t$$

Relax (t, z) with weight -2:

$$d(z) = \min(\infty, d(t) + (-2)) = 6 - 2 = 4, \pi(z) = t$$

Relax (y, t) with weight 8:

No update since $d(t) = 6$ and $d(y) + 8 = 7 + 8 = 15$ (larger)

Relax (y, x) with weight -3:

$$d(x) = \min(11, d(y) + (-3)) = 7 - 3 = 4, \pi(x) = y$$

Relax (y, z) with weight 9:

No update since $d(z) = 4$ and $d(y) + 9 = 7 + 9 = 16$ (larger)

Relax (x, z) with weight -4:

$$d(z) = \min(4, d(x) + (-4)) = 4 - 4 = 0, \pi(z) = x$$

Table after iteration 1:

Vertex	d	π
s	0	NIL
t	6	s
x	4	y
y	7	s
z	0	x

Iteration 2 (Relax all edges)

After performing all relaxations again, there are **no updates**, meaning the shortest distances are found.

Final Result

Vertex	Shortest Distance d	Predecessor π
s	0	NIL
t	6	s
x	4	y
y	7	s
z	0	x

The Bellman-Ford algorithm successfully finds the shortest paths from s. If an extra iteration results in changes, it means there is a negative-weight cycle, which is not the case here.

5. Calculate and find that $\Omega(n \log n)$ is a lower bound for computing Voronoi diagrams by reducing the sorting problem to the problem of computing Voronoi diagrams.

Step 1: Understanding Voronoi Diagrams

A **Voronoi diagram** partitions a plane into regions based on a given set of **n points**. Each region corresponds to a **generator point** such that any location in a region is closer to its generator than to any other point.

Step 2: Establishing a Lower Bound

To determine a **lower bound** on the complexity of computing Voronoi diagrams, we show that sorting **can be reduced** to computing a Voronoi diagram. Since sorting has a known lower bound of $\Omega(n \log n)$ in the comparison model, this implies that computing Voronoi diagrams cannot be done in less than $\Omega(n \log n)$ time.

Step 3: Reduction from Sorting to Voronoi Diagram Construction

To reduce sorting to Voronoi diagram construction, we follow these steps:

1. Given an unsorted set of numbers $S = \{s_1, s_2, \dots, s_n\}$, transform them into points in 2D space:

$$P = \{(s_1, 0), (s_2, 0), \dots, (s_n, 0)\}$$

This means all points are placed along the **x-axis**.

2. Construct the Voronoi Diagram for these n points.
3. Extract the sorted order of the original numbers from the Voronoi diagram:
 - Each Voronoi cell is defined by the perpendicular bisectors between points.
 - The order of the Voronoi regions along the **x-axis** directly gives the sorted order of S .
4. Since sorting requires $\Omega(n \log n)$ comparisons, and computing a Voronoi diagram has already sorted the points **implicitly**, it must also require at least $\Omega(n \log n)$ time.

Step 4: Conclusion

- Since sorting can be reduced to computing Voronoi diagrams in **O(1) additional time**, this means that computing Voronoi diagrams must be **at least as hard as sorting**.
- Sorting has a known lower bound of $\Omega(n \log n)$.
- Therefore, computing Voronoi diagrams has a lower bound of $\Omega(n \log n)$.

Thus, the reduction proves that **$\Omega(n \log n)$ is a lower bound** for computing Voronoi diagrams.

6. Professor Charon has a set of n sticks, which are piled up in some configuration.

Each stick is specified by its endpoints, and each endpoint is an ordered triple giving its $x; y; z$ coordinates. No stick is vertical. He wishes to pick up all the sticks, one at a time, subject to the condition that he may pick up a stick only if there is no other stick on top of it. Give a procedure that takes two sticks a and b and reports whether a is above, below, or unrelated to b.

Procedure to Determine Stick Order (Above, Below, or Unrelated)

Given two sticks aaa and bbb, each specified by their endpoints (x_1, y_1, z_1) and (x_2, y_2, z_2) , we need to determine whether:

1. Stick a is above stick b \rightarrow a must be removed before b.
2. Stick a is below stick b \rightarrow b must be removed before a.
3. Sticks a and b are unrelated \rightarrow They do not obstruct each other.

Step-by-Step Algorithm

To compare two sticks a and b, follow these steps:

1. Check if the sticks overlap in the (x, y) plane:
 - Project both sticks onto the xy-plane by ignoring the z-coordinate.
 - If the projections do not intersect, the sticks are unrelated.
2. Find the intersection point (if any) in the xy-plane:
 - Represent stick a and stick b as line segments in 2D.
 - Solve for the intersection point (x^*, y^*) of these segments.
3. Compare their heights at the intersection point:
 - Use linear interpolation to find the z-coordinate of each stick at (x^*, y^*) .
 - If $z_a > z_b$ then stick a is above stick b.
 - If $z_a < z_b$ then stick a is below stick b.
 - If $z_a = z_b$ they touch at that point but are still unrelated in stacking order.

Implementation of Above-Below Check

Here's a Python function to determine the relationship between two sticks:

Final Thoughts

- This approach ensures that **sticks are picked up in the correct order**.
- It runs in **O(1)** time per pair of sticks.
- It avoids **vertical sticks**, ensuring a valid comparison.

Explanation of the Code

1. `line_intersection(A1, A2, B1, B2) :`
 - Finds the intersection point of two line segments in the xy-plane.
 - Uses vector cross products to determine if they intersect.
2. `height_at_x(x, P1, P2) :`
 - Uses linear interpolation to find the **z-coordinate** of a stick at a given *x*-coordinate.
3. `stick_relation(A1, A2, B1, B2) :`
 - Calls `line_intersection` to check if the sticks overlap in xy-plane.
 - If they do, it finds their relative height at the intersection.
 - Returns whether **A is above B, below B, or unrelated**.