

```
# =====
# Assignment No. 8 (Group C - ML)
# Title: Binary Classification using Logistic Regression
# Dataset: Rain in Australia (Kaggle)
# =====
```

```
# Step 1: Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_curve, auc
```

```
# Step 2: Load Dataset (from official Rattle source)
url = "http://rattle.togaware.com/weatherAUS.csv"
data = pd.read_csv(url)
```

```
print("✅ Dataset Loaded Successfully!")
print("Shape:", data.shape)
print(data.head())
```

```
✅ Dataset Loaded Successfully!
Shape: (266946, 24)
   Date Location MinTemp MaxTemp Rainfall Evaporation Sunshine \
0  2008-12-01 Albury    13.4    22.9      0.6        NaN      NaN
1  2008-12-02 Albury     7.4    25.1      0.0        NaN      NaN
2  2008-12-03 Albury    12.9    25.7      0.0        NaN      NaN
3  2008-12-04 Albury     9.2    28.0      0.0        NaN      NaN
4  2008-12-05 Albury    17.5    32.3      1.0        NaN      NaN

   WindGustDir WindGustSpeed WindDir9am ... Humidity3pm Pressure9am \
0           W          44.0       W ...        22.0      1007.7
1          WNW          44.0      NNW ...        25.0      1010.6
2          WSW          46.0       W ...        30.0      1007.6
3            NE          24.0       SE ...        16.0      1017.6
4             W          41.0      ENE ...        33.0      1010.8

   Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm RainToday RISK_MM \
0      1007.1      8.0      NaN     16.9     21.8      No      0.0
1      1007.8      NaN      NaN     17.2     24.3      No      0.0
2      1008.7      NaN      2.0     21.0     23.2      No      0.0
3      1012.8      NaN      NaN     18.1     26.5      No      1.0
4      1006.0      7.0      8.0     17.8     29.7      No      0.2

   RainTomorrow
0            No
1            No
2            No
3            No
4            No

[5 rows x 24 columns]
```

```
# Step 3: Data Preprocessing
# Drop rows with too many missing values
data = data.dropna(subset=["RainToday", "RainTomorrow", "Humidity3pm", "Rainfall", "MaxTemp", "MinTemp", "WindGustSpeed"])

# Select important features
features = ["MinTemp", "MaxTemp", "Rainfall", "WindGustSpeed", "Humidity3pm", "RainToday"]
target = "RainTomorrow"

X = data[features]
y = data[target]

# Encode categorical features
le = LabelEncoder()
X["RainToday"] = le.fit_transform(X["RainToday"])
y = le.fit_transform(y) # Yes=1, No=0

# Handle missing values (if any)
X = X.fillna(X.mean())

/tmp/ipython-input-2120583147.py:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view
X[ "RainToday" ] = le.fit_transform(X[ "RainToday" ])
```

```
# Step 4: Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 5: Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Step 6: Build Logistic Regression Model
model = LogisticRegression(max_iter=200)
model.fit(X_train_scaled, y_train)
```

▼ LogisticRegression ⓘ ⓘ  
LogisticRegression(max\_iter=200)


```
# Step 7: Predictions
y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[:,1]
```

```
# Step 8: Evaluation
print("\n🎯 Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

🎯 Accuracy: 0.8465503401073133


		precision	recall	f1-score	support
	0	0.87	0.95	0.91	37158
	1	0.72	0.47	0.57	10180
accuracy				0.85	47338
macro avg		0.79	0.71	0.74	47338
weighted avg		0.84	0.85	0.83	47338

```
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
# Step 9: ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
```

```
plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.2f}")
plt.plot([0,1], [0,1], linestyle='--', color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Rain Tomorrow Prediction")
plt.legend()
plt.show()
```



```
# Step 10: Example Prediction
example = np.array([[15.0, 28.0, 2.0, 35.0, 60.0, 1]]) # RainToday = 1(Yes)
example_scaled = scaler.transform(example)
pred = model.predict(example_scaled)[0]

if pred == 1:
    print("☔ Prediction: It will RAIN tomorrow.")
else:
    print("☀️ Prediction: It will NOT rain tomorrow.")

☔ Prediction: It will NOT rain tomorrow.
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names,
warnings.warn(
```